# Into the Cacophony – Classification of Tweets on Twitter

David Sanford

*DSI-SM03*

*General Assembly*

(Dated: March 2017)

In this project I considered a classification scheme for tweets on twitter. I considered a data set of forty thousand tweets, evenly split between unfiltered tweets and a sample generated from a keyword filter using names of video game franchises. I trained a set of machine learning estimators using half the data set, resulting in cross-validation accuracies of between 75% and 83% on the training set and 76% for the best model on the validation set. I also outline the possible methods for organizing topic-related tweets into clusters based on genre or topic.

# I. DATA SET

Data for this project was streamed from twitter using the twitter API through the Tweepy library as an interfact in python. Tweets were loaded into a Mongo database running on a Docker container. Tweets loading was performed in the notebook Tweet-Loading.ipynb. Around three hundred thousand tweets were collected, but only two hundred thousand were in english and a fraction of those were used due to processing power constraints. The tweets were collected in the window of March 1-4, 2017.

For analysis, python was used with data held in the Pandas library. Pre-processing and machine learning were performed using Scikit-Learn. Data was pulled from the Mongo database through a spark intermediary database using a Pyspark interface. This was necessary due to memory constraints relative to the size of the database, from which only a small subset of the data was desired.

In terms of topic selection for classification, there are two apparent options for generating a large set of topic-related tweets. The first is to generate a curated list of users who tweets exclusively about that topic. This approach is difficult, as most users with a large volume of tweets will write off-topic tweets, which will contaminate the data set. Moreover, if the number of users chosen is too small the word choice of the individual users themselves will also be baked into the model, reducing classification effectiveness for unfiltered tweets.

The method used here is to gather topic-related tweets based on a series of keywords. For this approach to be viable, the topic chosen must have a sufficient number of keywords to produce a clean data set. Several topics were considered for this project before video games, including academic subjects, tabletop role-playing games, and tabletop board games. However, none of these subjects had sufficiently clean sets of keywords on which to filter tweets. Most video game franchises, however, are unique words or combinations of words, allowing for a large, relatively clean data set to be produced with a wide variety keywords.

The list of 86 keywords below was used to filter topic-related tweets.

Zelda, Tetris, Mario, Chrono Trigger, Street Fighter, Final Fantasy, Metroid, Half-Life, Resident Evil, Metal Gear, Castlevania, Pokemon, BioShock, SoulCalibur, StarCraft, Shadow of the Colossus, Doom, Diablo, World of Warcraft, Donkey Kong, Pac-Man, Halo, Deus Ex, Space Invaders, Sonic, Counter-Strike, Grim Fandango, Portal, Mass Effect, Last of Us, Star Fox, Mega Man, EarthBound, Prince of Persia, Call of Duty, Dark Souls, Perfect Dark, Ico, The Elder Scrolls, Skyrim, Morrowind, Silent Hill, Shenmue, Grand Theft Auto, Okami, Double Dragon, Red Dead, Galaga, Tomb Raider, Fallout, Uncharted, Assassin's Creed, Minecraft, Kingdom Hearts, Xenogears, Overwatch, Wii Sports, Wii Fit, The Sims, Terraria, Brain Age, Need for Speed, Lemmings, Madden NFL, Star Wars: Battlefront, Tom Clancy's, Duck Hunt, Splatoon, Super Smash, Dynasty Warriors, Monster Hunter, Kirby, Fire Emblem, Animal Crossing, God of War, Tekken, Garry's Mod, Myst, Angry Birds, Candy Crush, Fruit Ninja, Block Breaker, Doodle Jump, Space Invaders, Galaxian, Mortal Combat, Pong, Crysis

The list was taken from the Wikipedia lists of well-rated and high-selling games. Only key franchise words were used, and many older game names were dropped due to likely irrelevance. Franchise names that have broader contexts, for instance "Civilization" and "Battlefield" were also excluded to avoid contaminating the sample.

## II. DATA CLEANING

Data cleaning is performed in the notebook Initial_Loading_and_Cleaning.ipynb, using methods from lib/tweet_cleaning.py. Tweets hashtags, username references, retweet tags, emojis, links, and garbage strings containing letter/number combinations are all removed. Male and female proper names and surnames are also removed if they are within th 500 most popular names of their categoy based on 1990 U.S. Census data.

Strings from the keyword list are also removed based after the video game tweets are categorized by keyword. Only tweets containing a single keyword are included in the "video game" data set, and only tweets with no keywords are included in the "unfiltered" data set. From these data sets, twenty thousand tweets from each are chosen and divided evenly to produce equally-weighted training and validation sets of twenty thousand tweets apiece.

## III. LSA PRE-PROCESSING FOR BINARY CLASSIFICATION

The cleaned training set from Sec. II must be processed before it can be passed to classification methods. For this purpose I used a Latent Semantic Analsys (LSA) pipeline. The relevant code is included in the notebook Binary_Classification_Model.ipynb, using methods from the library

lib/lsa.py. The cleaned text is first vectorized using TF-IDF vectorization, using words appearing in between 0.1% (100) and 50% (5000) documents. Stop words were left in the sample, as removing them resulted in over 10% of the data set being reduced to empty sparse vectors. With stop words included, less than 0.5% of the training and validation sets were reduced to empty vectors.

The TF-IDF transformation resulted in 1172 features in sparse matrices, which were then fed through a truncated SVD transformation. The number of components kept was 200, which explained 56% of the total variance, as shown in Fig. 1. While not optimal, I judged this to be sufficient given the constraints in processing power.

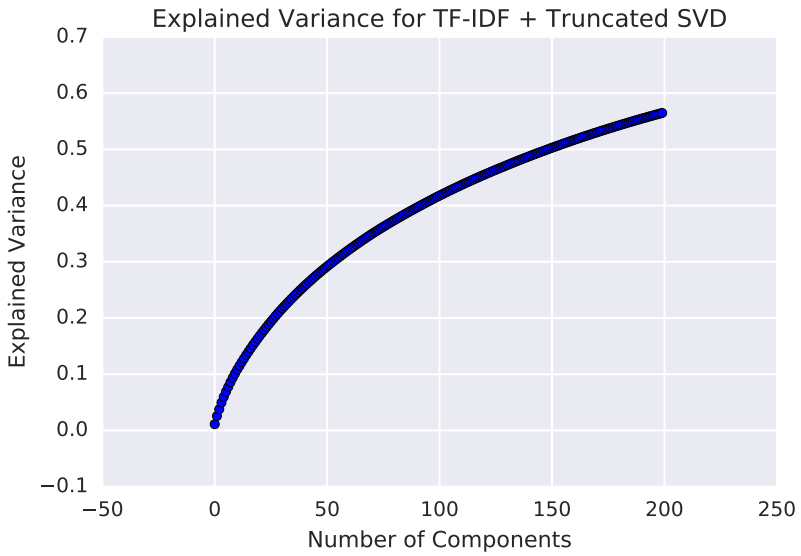


FIG. 1. *Figure 1: Explained Variance.* Total explained variance by number of components included for TF-IDF + Truncated SVD

## IV. BINARY CLASSIFICATION

Using the data sets generated through LSA in Sec. III, binary classification was performed using the six classification estimators in Table IV. Due to processing limitation, only the scikit-learn default parameters were used. Training was performed on 70% of the training set, using 5-fold cross-validation. A test accuracy was calculated using the remaining 30% of the training set.

The training and test accuracies for each estimator is shown in Table IV. The performance was reasonably similar for all the estimators used and almost identical for the training and test scores, ranging from 75% to 83% accuracy. This suggests that the correlations in the data are mostly linear, with variance resulting from "bleeding" of the classes into one another rather than number

| Estimator | Parameters |
|---|---|
| K-Nearest Neighbors | 5 neighbors, weighted by distance |
| Logistic Regression | L2 regression with C = 1 |
| SVC | RBF kernel, C = 1 |
| Decision Tree | Gini impurity |
| Random Forest | Gini impurity, 10 estimators |
| Extra Trees | Gini impurity, 10 estimators |

TABLE I. *Table 1: Estimators and parameters.* All parameters, both present and omitted, correspond to the scikit-learn defaults.

| Estimator | Train Accuracy | Test Accuracy |
|---|---|---|
| K-Nearest Neighbors | 0.783333 | 0.792444 |
| Logistic Regression | 0.774190 | 0.781444 |
| SVC | 0.822619 | 0.832889 |
| Decision Tree | 0.744524 | 0.751778 |
| Random Forest | 0.783048 | 0.792778 |
| Extra Trees | 0.790571 | 0.794000 |

TABLE II. *Table 2: Training accuracies.* Training accuracy is the average over 5-fold cross-validation. Test accuracy corresponds to the 30% of the training set not used in cross-validation.

of dimensions rather than meaninful non-linear correlations.

This is supported by the results of a random forest grid scan over the parameters given in Table IV, which showed only marginal improvement in performance. The random forest model was chosen for the grid search over the moderately better-performing SVC model due to a shorter training time.

Using the SVC estimator, the metrics in Table IV were calculated on the validation set. The

| Parameter | Values |
|---|---|
| 'criterion' | ['gini','entropy'] |
| 'n_estimators' | [2,5,10,20] |
| 'max_features' | [None, 'auto', 'sqrt','log2'] |
| Train Accuracy | 0.805905 |
| Test Accuracy | 0.813 |

TABLE III. *Table 3: Random forest grid search parameters and scores.* Training and testing accuracies are drawn from the same sets as in Table IV

| Model | Accuracy | Precision | Recall | F1-Score |
|---|---|---|---|---|
| SVC | 0.76 | 0.90 | 0.60 | 0.72 |

TABLE IV. *Table : Metrics on validation set.* All metrics used the entire validation set with the appropriate transformation applied to the raw data.

accuracy was moderately lower than the training accuracy, but still reasonably at 76%. More interestingly, the precision had a relatively high value of 90%, indicating that the identified sample of video game related tweets should be relatively clean. Recall, however, is only 0.6, indicating that only 60% video game related tweets in an unfiltered stream will be tagged.

This effect is even more pronounce in the ROC curve shown in Fig. 2. The false positive rate falls to less than 10% with the true positive rate still over 40%, so the model can easily be optimized for a low false positive rate while retaining a non-negligible true positive rate. However, achieving a high true positive rate requires acceptance of a significant false positive rate. A true positive rate of 80% requires a false positive rate of 40%, and a true positive rate of 90% requires a false positive rate of 90%.

## V. CLASSIFICATION MODEL PERFORMANCE AND IMPROVEMENT

The binary classification model examined here has a reasonable accuracy of around 75%, at least on the validation set. Further validation using a separate set of brand keywords is advisable, as the estimator may be overfit on brand-specific terms, and its true peformance for video games of all stripes may be lower. Consideration of a hand-selected set of tweets with manual classification for ultimate validation may also be desireable. On the positive side, this methodology should function for any topic with sufficiently distinct keywords, and it would be interesting to apply it to a corpus fo tweets selected using one of the terms already identified as ambiguous to judge performance out of the box.

Real gains in performance, however, likely require more careful curation of the origin data set and iterations on the analysis framework. Beyond the first step of accessing more processing power, some method of bootstrapping might provide better results. Many tweets are short, or even empty after keywords and stop words are removed, making them non-ideal for training. A procedure of concatenating random groups of the strings to artificially generate longer individual documents might improve performance. Alternatively, based upon the ultimate use case, a simple cut to remove excessively short and low-information tweets might be appropriate.

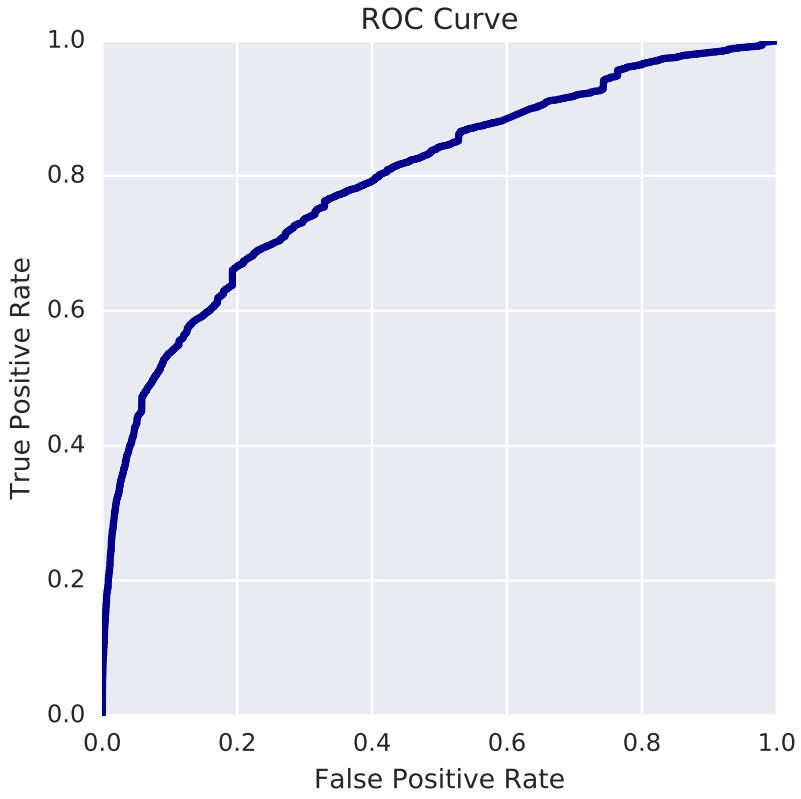


FIG. 2. *Figure 2: ROC Curve.* Results of variation in the threshold for the best-performing SVC model on the validation set.

Balancing the sub-classes and genres might also improve results. As shown in Table V, the relative appearance rate of various brands in the data set very weighted, with Zelda in particular appearing three times as much as the next most prevalent term (probably due to the time period in which the tweets were collected). Beyond this, balancing the number of tweets in the training sample by game genres might improve general results. That said, it is conceivable that some sub-topics are simply easier to classify than others – fantasy and science fiction games, for example, have more specialized vocabulary associated with them than those set in the real world.

## VI. GENRE/TOPIC CLUSTERING

Building upon binary classification, it would be desireable to classify tweets within sub-categories once they are identified as video game tweets. This could be an attempt to identify as related to a particular brand, to identify the genre of the game being discussed, or identify the topics being addressed in the tweet. The first use case likely reduces to keyword searches on

| Keyword | Percentage | Total |
|---|---|---|
| Zelda | 35.056 | 8764 |
| Overwatch | 11.116 | 2779 |
| Pokemon | 7.452 | 1863 |
| Minecraft | 5.260 | 1315 |
| Mario | 5.196 | 1299 |
| Halo | 3.072 | 768 |
| Sonic | 3.052 | 763 |
| Mass Effect | 2.980 | 745 |
| Resident Evil | 2.592 | 648 |
| Call of Duty | 2.108 | 527 |

TABLE V. *Prevalence of Keywords.* Top ten most prevalent keywords and the percentage of the training set that they compose.

brand-specific words, so the latter two possibilities are more interesting.

Neither avenue of inquiry has yet yielded results, so I will sketch possible approaches, based on code in Video_Game_Sub-Categorization.ipynb. Genre identification requires applying clustering algorithm to tweets processed through LSA, then attempting to identify genres based on the prevalence of various games within each cluster. Another approach might be to generate a corpus for each keyword out of all tweets associated with that keyword, then cluster those documents. Either of these approaches is hampered by the relative prevalence of tweets with particular keywords in the data set, as discussed above and shown in Table V. A more uniform selection of keywords through more careful curation is probably necessary for these purposes.

Topic modeling would instead focus on broad topics being discussed by twitter users, spanning different games and genres. Indentifying these topics would involve applying Latent Dirichlet Allocation (LDA) to the corpus and identifying key terms in each topic. An initial pass produced a grouping of naively unrelated terms. If topic modeling is appropriate for this corpus a significant amount of iteration in removing meaningless words would likely be necessary, and the technique may not be appropriate for the diversity of grammar present in tweets from a multitude of different users.

## VII. CONCLUSION

In this report I described a project to classify tweets on twitter by topic, using a sample drawn from a set of selected keywords over a large number of users. I fit a binary classification model which demonstrates 75% classification accuracy, and discussed possible methods to improve its performance in the future. I also discussed the initial stages and possible extension of two separate sub-classification schemes to extract further information from topic-related tweets.

The techniques described here, while simplistic, are broadly applicable. Going beyond keyword filtering to topic classification on a corpus is useful to discriminate any text on the internet, from comment threads and discussion boards to news stories and long blog posts. This identification could be used to select pages for an end user by a search engine, identify relevant documents for internal company use, or even select a clean set of documents for deeper NLP analyses.